

# Xiang Li

☎ (412) 600 2748    ✉ x16@andrew.cmu.edu    🔗 lxa9867.github.io

---

EDUCATION	<p><b>Ph.D. in Electrical and Computer Engineering</b>, Carnegie Mellon University 2021-2025 Supervised by <a href="#">Prof. Bhiksha Raj</a> (Full professor at Language Technologies Institute, SCS) Research Interests: Multimodal Perception / Generation</p> <p><b>M.S. in Electrical and Computer Engineering</b>, Carnegie Mellon University 2021-2022</p> <p><b>B.S. in Electrical Engineering</b>, Huazhong University of Science and Technology 2016-2020</p>
RESEARCH	<p>I am working on image/video generation with autoregressive/diffusion models and their corresponding autoencoder/tokenizer.</p>
EXPERIENCE	<p><b>Research Scientist, Google DeepMind</b>, Feb 2025-Now Veo post-training: One model with the joint capabilities of Camera, Reference, Motion, I2V, T2V, etc.. Vision tokenizer: Representation learning to find a joint space encoding RGB and controls.</p> <p><b>Research Intern, Adobe Research</b>, May 2024-Aug 2024 Supervised by Jason Kuen and Zhe Lin Research Topic: Autoregressive Visual Generation</p> <p><b>Research Intern, Microsoft</b>, Azure Computer Vision Team May 2023-Aug 2023 Supervised by Zicheng Liu, Yinpeng Chen, Chung-Ching Lin and Lijuan Wang Research Topic: Bridging Object Generation and Segmentation</p> <p><b>Research Intern, Tiktok</b>, Multimedia Lab May 2022-Aug 2022 Supervised by Shijie Zhao and Junlin Li Research Topic: Audiovisual Salient Object Detection</p> <p><b>Research Intern, Microsoft Research Asia</b> Oct 2020-Aug 2021 Supervised by Jinglu Wang Research Topic: Multimodal Video Instance/Object Segmentation</p>
PUBLICATIONS	<p>[<a href="#">ICML 2025</a>] Hao Chen, Yujin Han, Fangyi Chen, <b>Xiang Li</b>, Yidong Wang, Jindong Wang, Ze Wang, Zicheng Liu, Difan Zou, Bhiksha Raj “<i>Masked Autoencoders Are Effective Tokenizers for Diffusion Models</i>”</p> <p>[<a href="#">CVPR 2025</a>] Hao Chen, Ze Wang, <b>Xiang Li</b>, Ximeng Sun, Fangyi Chen, Jiang Liu, Jindong Wang, Bhiksha Raj, Zicheng Liu, Emad Barsoum “<i>SoftVQ-VAE: Efficient 1-Dimensional Continuous Tokenizer</i>”</p> <p>[<a href="#">ICLR 2025</a>] <b>Xiang Li</b>, Hao Chen, Kai Qiu, Jason Kuen, Jiuxiang Gu, Bhiksha Raj, Zhe Lin “<i>Image-Folder: Autoregressive Image Generation with Folded Tokens</i>”</p> <p>[<a href="#">ICLR 2025</a>] Xiaohao Xu, Tianyi Zhang, Shibo Zhao, <b>Xiang Li</b>, Sibow Wang, Yongqi Chen, Ye Li, Bhiksha Raj, Matthew Johnson-Roberson, Sebastian Scherer, Xiaonan Huang “<i>Scalable Benchmarking and Robust Learning for Noise-Free Ego-Motion and 3D Reconstruction from Noisy Video</i>”</p> <p>[<a href="#">Arxiv 2024</a>] Hao Chen, Abdul Waheed, <b>Xiang Li</b>, Yidong Wang, Jindong Wang, Bhiksha Raj, Marah Abdin “<i>On the Diversity of Synthetic Data and its Impact on Training Large Language Models</i>”</p> <p>[<a href="#">NeurIPS 2024</a>] Hao Chen, Yujin Han, Diganta Misra, <b>Xiang Li</b>, Kai Hu, Difan Zou, Masashi Sugiyama, Jindong Wang, Bhiksha Raj “<i>Slight Corruption in Pre-training Data Makes Better Diffusion Models</i>”</p> <p>[<a href="#">NeurIPS 2024</a>] Kai Hu, Weichen Yu, Tianjun Yao, <b>Xiang Li</b>, Wenhe Liu, Lijun Yu, Yining Li, Kai Chen, Zhiqiang Shen, Matt Fredrikson “<i>Efficient LLM Jailbreak via Adaptive Dense-to-sparse Constrained Optimization</i>”</p>

- [[NeurIPS 2024](#)] Hao Chen, Ankit Shah, Jindong Wang, Ran Tao, Yidong Wang, **Xiang Li**, Xing Xie, Masashi Sugiyama, Rita Singh, Bhiksha Raj “*Imprecise Label Learning: A Unified Framework for Learning with Various Imprecise Label Configurations*”
- [[ECCV 2024](#)] **Xiang Li**, Kai Qiu, Jinglu Wang, Xiaohao Xu, Rita Singh, Kashu Yamazaki, Hao Chen, Xiaonan Huang, Bhiksha Raj “*R<sup>2</sup>-Bench: Benchmarking the Robustness of Referring Perception Models under Perturbations*”
- [[CVPR 2024](#)] **Xiang Li**, Jinglu Wang, Xiaohao Xu, Rita Singh, Yan Lu, Bhiksha Raj “*Towards Robust Audiovisual Segmentation in Complex Environments with Quantization-based Semantic Decomposition*”
- [[MICCAI 2024](#)] Yizhou Zhao, Hengwei Bian, Michael Mu, Mostafa R Uddin, Zhenyang Li, **Xiang Li**, Tianyang Wang, Min Xu “*Training-free CryoET Tomogram Segmentation*”
- [[InterSpeech 2024](#)] Muqiao Yang, **Xiang Li**, Umberto Cappellazzo, Shinji Watanabe, Bhiksha Raj “*Evaluating and Improving Continual Learning in Spoken Language Understanding*”
- [[ICML 2024](#)] Hao Chen, Jindong Wang, Lei Feng, **Xiang Li**, Yidong Wang, Xing Xie, Masashi Sugiyama, Rita Singh, Bhiksha Raj “*A General Framework for Learning from Weak Supervision*”
- [[ICML 2024](#)] **Xiang Li**, Yinpeng Chen, Chung-Ching Lin, Rita Singh, Bhiksha Raj, Zicheng Liu “*Completing Visual Objects via Bridging Generation and Segmentation*”
- [[NAACL 2024](#)] Zhaorun Chen, Zhihong Zhu, Ruiqi Zhang, **Xiang Li**, Bhiksha Raj, Huaxiu Yao “*AutoPRM: Self-supervised Fine-grained Feedback for Multi-Step Reasoning via Controllable Question Decomposition*”
- [[ICASSP 2024](#)] Muqiao Yang, Umberto Cappellazzo, **Xiang Li**, Shinji Watanabe, Bhiksha Raj “*Improving Continual Learning of Acoustic Scene Classification via Mutual Information Optimization*”
- [[ICPR 2024](#)] Jimin Sohn, Haeji Jung, Zhiwen Yan, Vibha Masti, **Xiang Li**, Bhiksha Raj “*Fashion Image Retrieval with Occlusion*”
- [[Computer Speech & Language 2024](#)] Fan Yang, Muqiao Yang, **Xiang Li**, Yuxuan Wu, Zhiyuan Zhao, Bhiksha Raj, Rita Singh “*A Closer Look at Reinforcement Learning-based Automatic Speech Recognition*”
- [[EMNLP 2023](#)] **Xiang Li**, Jinglu Wang, Xiaohao Xu, Muqiao Yang, Rita Singh, Bhiksha Raj “*Towards Noise-Tolerant Speech-Referring Video Object Segmentation: Bridging Speech and Text*”
- [[NeurIPS 2023](#)] **Xiang Li**, Chung-Ching Lin, Yinpeng Chen, Jinglu Wang, Zicheng Liu, Bhiksha Raj “*PaintSeg: Training-free Segmentation via Painting*”
- [[ICCV 2023](#)] **Xiang Li**, Jinglu Wang, Xiaohao Xu, Xiao Li, Bhiksha Raj, Yan Lu “*Robust Referring Video Object Segmentation via Relational Multimodal Cycle Consistency*”
- [[ACM MM 2023](#)] **Xiang Li**, Yandong Wen, Muqiao Yang, Jinglu Wang, Rita Singh, Bhiksha Raj “*Rethinking Voice-Face Correlation: A Geometry View*”
- [[InterSpeech 2023](#)] Liao Qu\*, Xianwei Zou\*, **Xiang Li\***, Yandong Wen, Rita Singh, Bhiksha Raj “*The Hidden Dance of Phonemes and Visage: Unveiling the Enigmatic Link between Phonemes and Facial Features*”
- [[AAAI 2023](#)] **Xiang Li**, Haoyuan Wang, Shijie Zhao, Junlin Li, Li Zhang, Bhiksha Raj “*Panoramic Video Salient Object Detection with Ambisonic Audio Guidance*”
- [[AAAI 2022](#)] **Xiang Li**, Jinglu Wang, Xiao Li, Yan Lu “*Hybrid Instance-aware Temporal Fusion for Online Video Instance Segmentation*”

[[TMM 2022](#)] **Xiang Li**, Jinglu Wang, Xiao Li, Yan Lu “*Video Instance Segmentation by Instance Flow Assembly*”

[[Ubicomp 2022-HASCA](#)] **Xiang Li\***, Jinqi Luo\*, Rabih Younes “*ActivityGAN: Generative Adversarial Networks for Data Augmentation in Sensor-Based Human Activity Recognition*” (**Best Paper Award**)

[[TWC 2022](#)] Yong Huang, **Xiang Li**, Wei Wang, Tao Jiang, Qian Zhang “*Forgery Attack Detection in Surveillance Video Streams Using Wi-Fi Channel State Information*”

[[WCNC 2022](#)] Chenhui Zhao, **Xiang Li**, Rabih Younes “*Self-supervised Multi-Modal Video Forgery Attack Detection*”

[[INFOCOM 2021](#)] Yong Huang, **Xiang Li**, Wei Wang, Tao Jiang, Qian Zhang “*Towards Cross-Modal Forgery Detection and Localization on Live Surveillance Videos*”

[[IJCAI 2020-DL-HAR](#)] Jinqi Luo, **Xiang Li**, Rabih Younes “*Towards Synthetic Data Generation and Interpretation on Sensor-based Fine-grained Hand Activity Recognition*”

## TEACHING

### Graduate Teaching Assistant at CMU

- 11785, Introduction to Deep Learning, Fall 2021
- 18661, Introduction to Machine Learning, Fall 2023
- 18789, Deep Generative Modeling, Spring 2024

## ACADEMIC SERVICE

### Reviewer

- Conference: CVPR, ECCV, ICCV, ACL, EMNLP, NAACL, NeurIPS, ICML, ICLR
- Journal: TIP, TMM, TWC